Home My Network Jobs Messaging

Vincent Wang
Individual article

Manage Next

Style

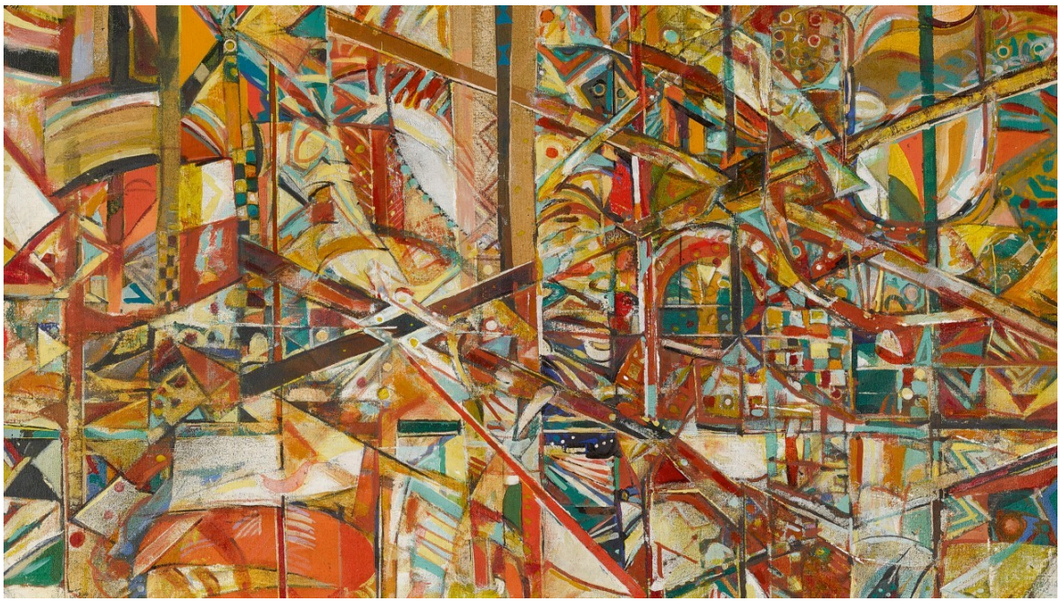

Add credit and caption

The Alignment Tax: What AI Knowledge Systems Must Solve

The meeting was scheduled for an hour. Eight people, three dialing in from regional offices. The agenda had four items. The last was the one that mattered: a resource allocation decision that would shape the next two quarters.

Twenty minutes in, someone from finance mentioned a number. It was not new. It had appeared in an email three weeks earlier, sent to a distribution list that included five of the eight people in the room. The other three - including the two regional leads whose input was the reason the meeting existed - had never received it.

The meeting paused. Someone shared the attachment. The finance lead explained the context. A clarifying question led to a second, then a third. Twenty minutes passed. The conversation was useful, but it was not the conversation anyone had come to have.

The meeting ended. The allocation decision was deferred. No one described what happened as a failure, because it felt like a normal meeting. It was a normal meeting. That is the problem.

Draft - saved

What occurred in that room was not a scheduling error or a communication lapse. It was a structural failure in how the organization distributes knowledge. It happens every day, in every company past a certain size, and it does not have a name. So give it one: the alignment tax.

The failure in that meeting had three layers, each with a different character and a different cost.

The first is **information misalignment**: the gap between the people who hold a document and the people who need it. The finance memo existed. It had been sent. The failure was in distribution, invisible to both sides. The sender could not have known who would need the number three weeks later. The people who needed it could not have known where to look, and so they did what everyone does: asked one colleague, then another, traded emails across departments, until the file surfaced from wherever it lived. This is the layer organizations recognize. They have thrown shared drives at it, enterprise search, Slack, Teams, producing vast collections of documents that are technically accessible and practically unfindable. The residual cost lands on the assistants and operations staff who compose every email as if for an archive, push files to anyone who might need them, and send "just in case" copies they know will go unread. The missing infrastructure becomes a quiet, continuous tax paid by the people least positioned to refuse it.

The second is **progress misalignment**. In any organization running dozens of parallel initiatives, maintaining a single accurate picture of current status is nearly impossible; not because people refuse to report, but because they report into different systems. Team A keeps raw data in handwritten notebooks. Team B tracks everything through a ticketing system. Team C has a meticulous file structure in OneDrive. Projects stall while the organization treats them as active. Projects finish while the budget and attention keep flowing. Progress information is generated continuously but recorded in scattered formats, captured at review cycles too infrequent to reflect what is actually happening.

The third layer almost never gets named. Call it **ignorance misalignment**, and it is blind on both sides. Those three regional leads could not have prepared for the finance data, because they did not know the data existed. But the organizer could not have known they were missing it either. She had the memo. So did the two colleagues she works with daily. From where she sat, the number was common knowledge. The gap was invisible to the people on both sides of it. You cannot ask for what you do not know to ask for, and you cannot share what you do not realize others lack. This is not a gap in expertise or judgment. It is a gap in the conditions that make expertise usable — the common floor of facts, definitions, and situational context that every substantive conversation assumes is already in place.

Here is what this looks like at scale. Last month, I ran a one-hour technical session for roughly a hundred colleagues, demonstrating how to build a custom email agent from scratch. The session had been carefully designed: a detailed prompt breakdown, a structured path for refining the agent in real time, and an introduction to applying it in other scenarios. Of the full hour, the core demonstration took twelve minutes. Most people

Draft - saved

eventually followed along. Since last week, they have been using the agent daily.

The remaining forty-eight minutes went to questions that had nothing to do with the idea itself. Where do I install this? Can you send the access link? Which folder does this file go in? How do I configure the connector? None of these people had failed to prepare. Most did not know there was anything to prepare for. And I could not have known what they were missing. To me, an AI agent demo obviously requires AI tools already installed. That assumption was so basic it never surfaced as something to verify. The curse of knowledge made the gap invisible from my side; the absence of knowledge made it invisible from theirs. And knowledge can be anything from "AI deployment and Vibe Coding" to "We have received numbers from Q3." The curse of knowledge in everything leads to misalignment of ignorance everywhere.

A hundred capable professionals, each blind to a different prerequisite, not from negligence, but because no part of the organization owned the job of surfacing what they didn't know. And what makes this layer expensive is what keeps it invisible. Most of those who could not install the tool did not raise their hands. Asking a baseline question in front of senior colleagues - when everyone else appears or pretends to follow - is a risk people correctly calculate and decline to take. In my demo, the material was accessible enough that everyone eventually reached the same floor, and we managed to teach the material effectively to almost everyone in 12 minutes. But anything more complicated, and the substance reaches only the fraction who were already there. And those who remain? Well, fake it till you make it.

Research from McKinsey and Bain converges on a striking number: nearly one working day in five, for every senior leader, is consumed by alignment work. Synchronizing context, rebuilding shared ground, catching people up. Before any substantive discussion begins. Too frequently, one day is still not enough.

This is not the cost of complexity. It is the cost of running a complex organization without infrastructure for its own cognition.

I study the microhistory of technology: the everyday devices and infrastructure that quietly restructure how people live. How flush toilets reshaped domestic architecture. How industrial bleach defined the white color of a wedding gown. My professional reflex, when I encounter any technology, is not just to ask what it can do but what the people around it adopt, resist, and transform with technology.

Once upon a time, when organizations faced a knowledge-distribution problem at this scale, the answer was the vertical filing cabinet. When the Library Bureau began promoting it in the 1890s, it was sold as a space-saving device. What it actually did was transfer institutional memory out of the heads of senior clerks and into a physical system that any employee could query. Companies discovered they could reassign a client

Draft - saved

relationship without a ritual handover, or dismiss a veteran employee without losing the business he carried in his head. An efficiency product turned out to be an epistemic one: better tools for managing, coordinating, and controlling internal information, which meant better products. The historian JoAnne Yates called the shift that followed "control through communication." The phrase is worth remembering, because the AI tools we are now deploying will be asked to do something of the same order.

The dominant responses to the alignment problem do not rise to that order. Some organizations funnel everything into a single company-wide knowledge base and hope retrieval will close the gaps. Others lean on a handful of experienced managers to carry the organization across the spaces the system cannot reach. Neither response is wrong. Both are incomplete in the same way: they flatten categories of knowledge that are not the same kind of thing, and then ask one mechanism to hold all of them.

What the three layers of misalignment actually need is three layers of knowledge base.

At the outermost ring sits the Corporate Knowledge Base - the organization's institutional memory. Archived projects, policy decisions, and closed initiatives. It belongs to the company. When a new employee joins, the portions relevant to their role are surfaced into their workspace, and an onboarding agent walks them through context interactively. The same induction a good team lead would run, but consistent, information-dense, and not dependent on whichever senior colleague happens to be free that week. When the new member meets the team, they get to meet the people and discuss the nuances instead of having to walk through the basics. When a new initiative launches, it draws context from here. When an initiative ends, its contents fold back in.

In the middle sits the Initiative Knowledge Base. Each active project has its own bounded knowledge space, shared among the people working on it and whoever needs to track its progress. This is where the living state of current work exists: what is active, what has stalled, what the team has learned this week. Management can access progress data whenever they need it. When the work is done, the base is closed, structured, and handed up to the Corporate layer.

At the innermost ring sits the Personal Knowledge Base. This space belongs to the individual for the duration of their employment. It holds what a person actually carries through their working life: their private judgments, their notes on relationships and clients, the instincts they have developed over years. No organizational agent reads it. No enterprise system indexes it. Information flows in freely; it flows out only when the person chooses to contribute something upward, into an initiative or a report.

This layer exists to clear the ground. A financial analyst whose real value lies in the conclusions they draw from a dataset currently spends hours building the spreadsheet, formatting the report, and translating a judgment she already holds into language that fits the document's

Draft - saved

structure. The formatting, the retrieval, the mechanical verbalization. . . . None of it requires them. It requires a person nowadays only because no alternative has existed. The Personal Knowledge Base, configured to their patterns, absorbs the mechanical substrate so their cognitive space is freed for the work that actually demands them: the analysis, the judgment, the conclusion no system can reach on their behalf.

One commitment makes this architecture trustworthy. The company never reads the Personal Knowledge Base. When an employee departs, the corporate-context portion, client preferences, internal dynamics, and proprietary workflows are deleted. The architecture does not pretend to own what lives in the person's head, and it does not let the person walk away with what belongs to the company. The principle is the precondition: without the guarantee of privacy, the other two layers lose the trust that makes them safe to use. If people fear the system is watching them rather than serving them, they input defensively and selectively — and the resulting data does more damage than no data at all.

Enterprise software has been circling these needs for years. Oracle, Microsoft, Salesforce, and dozens of others. The problem is not the absence of tools but the walls between them. Salesforce updates do not flow into email drafts. ClickUp submissions do not sync into Excel workbooks. The corporate and initiative layers live in proprietary shells, hard to bridge across platforms. The personal layer, meanwhile, has existed only in the person's head, forcing them to handle every mechanical step by hand, every time, in addition to serving as the bridge between those proprietary Apps. What has changed is that LLMs are now capable enough to navigate the compatibility gaps between data structures and software, smoothing the technical pathways across all three layers. Knowledge-base architectures like Karpathy named LLM-Wiki or Sigman and Jovovich named MemPalace produce storage and retrieval in plain-language markdown that anyone can read. Digitizing institutional knowledge is no longer the episodic privilege of engineers and data analysts. It is becoming the continuous convenience - and responsibility - of everyone. You take the convenience of maintaining your personal knowledge base, and you take responsibility for contributing to the team's and company's knowledge infrastructure.

The architecture respects what humans cannot cleanly express. When it falls short, it can try again and improve simply by being used more and more. It does not need a perfect setup on day one. It clears the mechanical grind and opens the ground for the work only a human can do. But it does not assume humans will use it perfectly. People will skim their briefings. Files will be ingested that the person never reads. Privacy and security will surface new problems, and much of the time saved must be reinvested in learning and tweaking in the earlier adoption period. LLMs and GenAI are not magic-blessed miracles. They are human-developed configurable tools. They can do a great deal, but we, individually and collectively, have to learn to work with them.

Draft - saved

Every organization is already paying the alignment tax. It shows up in the meeting that ran to catch up instead of making decisions, in the assistant who composes every email as if for an archive because the archive does not exist, and in the capable professional who did not raise a hand because the cost of raising it was higher than the cost of staying lost. In every instance of "fake it till you make it." The tax is continuous and expensive. Nothing in return for it is ever booked.

The question worth asking is not how to make people faster, or meetings shorter, or documents more findable. It is prior to all of those.

How does your organization deal with what it doesn't know?

And if you had to answer that question by the end of the afternoon, not next quarter, not after the next survey, who would you even ask?

Draft - saved